

Lightning Pool: A Non-Custodial Channel Lease Marketplace

Olaoluwa Osuntokun
roasbeef@lightning.engineering

Conner Fromknecht
conner@lightning.engineering

Wilmer Paulino
wilmer@lightning.engineering

Oliver Gugger
oliver@lightning.engineering

Johan Halseth
johan@lightning.engineering

October 28, 2020

Abstract

In this paper we examine the core inbound liquidity bootstrapping problem the Lightning Network faces, and frame it as a resource allocation problem to be solved by the application of market design and auction theory. We present Lightning Pool, a non-custodial channel lease marketplace, implemented as a sealed-bid frequent batched uniform clearing price auction which allows participants to buy and sell capital obligations on the network. We call these capital obligations *channel leases*. A channel lease can be viewed as a cross between a traditional fixed-income asset and an internet bandwidth peering agreement. Channel leases allow nodes on the network with idle capital to earn yield (based on a derived *per-block* interest rate) by selling channels to other agents in the marketplace. The duration of such contracts is enforced on-chain using Bitcoin Script. We construct **Lightning Pool** using a novel design for constructing overlay applications on top of Bitcoin: a shadow chain. Shadow chains allow users to remain in full custody of their funds at all times while also participating in higher level applications.

Contents

1	Introduction	4
1.1	Our Contributions	6
2	Preliminaries	6
3	Background	6
3.1	Payment Channels & the Lightning Network	7
3.2	Bootstrapping Problems in the Lightning Network	8
3.3	New Routing Node Bootstrapping	9
3.4	New Service Bootstrapping	9
3.5	End User Bootstrapping	11
3.6	Market Design & Auction Theory	11
3.7	Money Markets & Capital Leases	12
4	Bootstrapping Problems as Solved by CLM	13
4.1	Bootstrapping New Users via Sidecar Channels	13
4.2	Demand Fueled Routing Node Channel Selection	13
4.3	Bootstrapping New Services to Lightning	14
4.4	Cross-Chain Market Maker Liquidity Sourcing	14
4.5	Instant Lightning Wallet User On Boarding	15
4.6	Variance Reduction in Routing Node Revenue	15
5	The Channel Lease Marketplace	15
5.1	High-Level Description	15
5.2	Lightning Channel Leases	18
5.3	Non-Custodial Auction Accounts	19
5.4	Order Structure & Verification	20
5.5	Auction Design	22
5.5.1	Auction Specification	22
6	The Shadowchain: A Bitcoin Overlay Application Framework	24
6.1	High-Level Description	25
6.2	Comparison To Related Frameworks	26
6.3	The Shadowchain Framework	26
6.3.1	Shadowchain Orchestrator	26
6.3.2	Lifted UTXOs	26
6.3.3	The Shadowchain	27
6.3.4	Shadowchain Operation	30
7	Lightning Pool: A Channel Liquidity Marketplace as a Shadow Chain	31
7.1	Instantiating a CLM	32
7.1.1	System Initialization	32
7.1.2	Lightning Pool Accounts	33
7.1.3	Channel Leases in the Lightning Network	36
7.1.4	Order Structure	38
7.1.5	Node Rating Agencies	39
7.1.6	Uniform Price Market Clearing & Matching	40
7.1.7	The Batch Execution Transaction	41

7.2	The Lightning Pool Shadowchain	43
7.2.1	Lightning Pool Accounts as Lifted UTXOs	43
7.2.2	Auction Batch Proposal	43
7.2.3	Shadowchain Batch Execution	44
7.2.4	Unconfirmed Batch Cut-Through	44
7.2.5	Auction Upgrades	44
8	Security Analysis	45
9	Future Directions	45
10	Related Work	46
11	Conclusion	46
12	Acknowledgments	46

1 Introduction

Lightning Network Bootstrapping Challenges. The Lightning Network (LN) (cite) is the largest deployed Layer 2 payment channel network (cite <https://lightning.network/lightning-network-paper.pdf>). A payment channel network is comprised of a series of individual payment channels, which, when strung together, enable rapid low-latency payments between participants in the network. Due to the off-chain nature of these payments (only the final summary hits the blockchain), the cost of payments on the LN are typically much lower than an equivalent payment on the base blockchain (cite). In order to be able to send funds on a network, a user must open a payment channel (cite) to another participant on the network. Once the channel has been opened, both participants are able to send and receive a nearly unbounded number of payments off-chain, possibly never closing the channel on-chain. Similarly, in order to receive on the network, a user requires *another individual* to open a channel *to* the receiver. A participant can only send and receive up to the total amount of Bitcoin in a channel committed to by both parties. This allocation of capital by one party to enable another party to receive funds on the network is typically referred to as *inbound liquidity* or *inbound bandwidth*. Because a would-be user of the network must somehow convince *others* to allocate capital towards them in order to receive funds on the network, the inbound bandwidth problem remains a significant barrier to the adoption and bootstrapping of the Lightning Network.

Routing Node Capital Requirements. In order to incentivize users to commit capital to the network so as to help other users of the network send and receive payments, each time a node forwards a payment successfully they receive a fee. As commonly implemented, this fee has a fixed base amount to be paid for all forwards independent of payment size, and a fee rate or proportional amount that must be paid for each millionth of a satoshi forwarded (cite). We refer to nodes that join in order to facilitate payments and collect forwarding fees as *routing nodes*. In order to forward a payment of size N on the network, a routing node needs to have $N_{in} - F$ Bitcoin allocated *towards* it as inbound bandwidth, and N_{out} Bitcoin allocated to another node as outbound bandwidth. The factor of F is the fee collected by the routing node, with the following constraint being met $F = N_{out} - N_{in}$. Due to this requirement that there be sufficient inbound *and* outbound bandwidth, even those nodes that exist solely to help other nodes send and receive payments, themselves face the inbound bandwidth allocation problem, a fundamental bootstrapping issue for the LN.

Market Design Resource Allocation Problems. The field of market design is a sub-field of economics which is concerned primarily with the efficient allocation of scarce resources (cite). Within this sub-field, of interest is a branch of market design concerned with instances wherein money is used to govern the exchange of goods and services: auction design (cite). Auction design can be used to effectively allocate scarce resources within a domain. Common established examples of market design widely used today include: carbon emissions credits, electricity markets, auctions for airport gates, and wireless spectrum auctions (cite) (cite) (cite). In each of these examples, market design is used to allow more effective communication of pricing information, resource availability, and the expression of preferences (cite). Our first insight is framing the solution to the inbound bandwidth bootstrapping problem within the lens of market de-

sign. In the context of the LN, the scarce resource we aim to more efficiently allocate is inbound channel bandwidth.

LN Bootstrapping as Resource Allocation Problem. In the absence of a proper venue, those that need inbound capital to operate their Lightning services are forced to solicit capital on various chat groups, forums, or public venues like Twitter. On the other side, those seeking to deploy capital in order to facilitate network operation and gain routing fees must guess as to exactly where their capital is most demanded. As node operators may not necessarily know where their capital is most demanded, they risk opening channels to locations where they aren't actually need, leading to poor resource utilization and capital inefficiency. It's as if node operators are speculatively building roads that no one will use (why isn't my node forwarding?), and those seeking to receive aren't able to flag their service as an attractive destination to be connected to internal network "highways". Lightning Pool solves this resource allocation problem using an auction that matches up those seeking to deploy capital (by opening channels) to those that need these channels to operate their Lightning service or business. With each executed batch, the participants of the auction derive a *per block* interest rate which is effectively the current lease rate for capital on the Lightning Network. The auctioneer or an independent agency is also able to provide Node Ratings to participants of the network, which can be used to make more informed decisions with respect to the *quality* of the channel lease being purchased. (cite predatory routing hi-jacking stuff)

Channel Lease Marketplaces. In this paper, we present Lightning Pool, a non-custodial channel lease marketplace that draws on modern auction theory to construct an auction that enables participants to buy and sell inbound channel bandwidth. Participants in the marketplace buy and sell a channel capital obligation, which we call a Lightning Channel Lease (LCL). An LCL is similar to a traditional bond in that one party acquires capital from another for productive use, with the party parting with their capital being compensated for their cost of capital. However as the funds within an LCL can only be used in the Lightning Network for sending/receiving, an LCL is analogous to the creation of a new virtual "road" within the LN connecting two destinations. Critically, when one purchases an LCL, the period of time those funds must be committed is enforced on-chain using Bitcoin Script. As a result, buyers of inbound channel bandwidth can be sure the capital will be committed for a set period of time. The auction itself contains several sub-auctions for the exchange of particular duration intervals expressed in blocks (similar to the various U.S Treasury auctions (cite)). A non-trusted auctioneer facilitates the marketplace by accepting sealed-bid orders, clearing the market using a uniform clearing price for each duration bucket, and finally executing batches of contracts using execution transactions that update all involved accounts, delivering the purchased channels to all parties in an atomic manner. Lightning Pool and the LCL solve the inbound bandwidth problem by allowing participants on the network to effectively exchange pricing signals to determine exactly *where* in the network capital should be allocated.

Shadow Chains as an Application Framework. Lightning Pool is the first application built on top of Bitcoin that utilizes the *Shadow Chain* paradigm to construct an application-specific overlay system on top of existing Bitcoin unspent transaction outputs (UTXOs). A user joins a shadow chain by creating a special multi-signature based output using a public key provided by the target

shadow chain manager. Once a user has joined the shadow chain, proposed state transitions are packaged up (in the form of a block) by the shadow chain manager and proposed to each active participant. A shadow chain block updates the application state of all participants and is embedded within a normal Bitcoin transaction. Depending on the application, a shadow chain block may be indistinguishable from a normal Bitcoin transaction. Shadow chains are able to *compress* state transitions *off-chain* by employing a form of multi-party transaction cut-through (cite). As participants of a shadow-chain remain in custody of their funds at all times, complex fraud proofs or exit games are unnecessary, significantly simplifying the implementation of a given shadow chain. We note that the shadow chain concept itself can be used for other applications as well.

1.1 Our Contributions

In summary, we make the following contributions:

- We propose a solution to the inbound bandwidth problem of the LN in the form of a marketplace to buy and sell inbound bandwidth obligations.
- We propose the Lightning Channel Lease, an inbound bandwidth obligation contract that pays a per-block interest rate to the seller from the buyer, and whose duration is enforced on-chain with Bitcoin Script.
- We put forth the concept of a Node Rating agency for channel leases in order to provide marketplace participants with information about the *quality* of a channel lease.
- We construct a new system, Lightning Pool, which is a non-custodial marketplace with off-chain order submission and on-chain batch execution that allows parties to exchange LCL contracts in an atomic manner.
- We design a new Bitcoin application design framework, the shadow chain, which can also serve other applications.

Organization. Fill in final organization.

2 Preliminaries

not needed?

3 Background

In this section, we aim to introduce some necessary background that will be built upon in later chapters to construct our solution. First, we'll describe multi-hop payment channels and the Lightning Network as deployed today. Next, we'll explore the nature of the inbound bandwidth bootstrapping problem the Lightning Network faces today. Along the way, we'll explain the dynamics of routing nodes in the network, as they're a key component of the system. Next, we introduce the field of market design and more specifically the sub-field of auction design, to demonstrate how auction design can be used to solve resource allocation problems in the real world. Next, we provide some brief background

on money markets in the traditional financial system, and how this relates to our concept of channel leases.

3.1 Payment Channels & the Lightning Network

Basic Payment Channels

A payment channel (cite) in its simplest form (cite) is an on-chain 2-of-2 multi-signature output created by parties A and B . One or both parties deposit funds into a Bitcoin Script output constructed using two public keys P_a and P_b . The transaction that creates this multi-sig output is referred to as the *funding transaction*. Before broadcasting the funding transaction, another transaction dubbed the *commitment transaction* is constructed using a series of agreed upon parameters by the two parties (cite bolt). The commitment transaction spends the funding transaction and creates two new outputs D_a and D_b that if broadcast, will *deliver* the up-to-date balances allocated between the parties to the channel. Once the funding transaction is confirmed and broadcasted, both parties are able to rapidly update the balance of the delivery outputs, D_a and D_b so as to facilitate efficient payments between the parties.

Bi-Directional Payment Channels

In order to safely make bi-directional payments between both parties to a payment channel, modern channel designs also employ a *commitment invalidation mechanism* (cite paddy) to ensure that only the latest commitment transaction state can be broadcasted and redeemed via the underlying blockchain. The most commonly used commitment invalidation scheme is the **replace-by-revocation** construct. In this construction, during channel negotiation, a security parameter T (which may be asymmetric for both parties) is negotiated. Using this value T which is typically expressed in blocks, a commitment transaction state can only be fully redeemed by the broadcasting party after a period of T blocks has passed. During this interval, the non-broadcasting party $P_{defender}$ is able to provide the contested delivery output D_{a_i} with a valid witness W_{r_n} which *proves* that there exists a *newer* state n with $n > i$ that has been ratified by both parties. The exact details of this construct are outside the scope of this paper, but Bitcoin Script and basic cryptography are used to allow a defending party to present an objective statement of contract violation by the opposing party.

Hash Time Lock Contracts & Multi-Hop Payments

The final component of modern multi-hop payment channels is the Hash Time Locked Contract, or the HTLC. The HTLC enables payments to travel over a *series* of payment channels, allowing a set of interlinked payment channels to be composed into a logical *payment network*. An HTLC can be viewed as a specific case of a time locked commit and reveal puzzle (cite ranjit?). Loosely, an HTLC consists of four parameters: the public key of the sender P_s , the public key of the receiver P_r , the payment amount expressed in satoshis (cite) A_{sat} , a payment secret r s.t. $H(r) = h$, and an absolute block timeout T . Given these

parameters, a Bitcoin Script is set up such that, the funds deposited in the script hash output can be redeemed by the receiver P_r via a public key signature by their public key and the revelation of the payment pre-image r , or by the sender P_s after the absolute timeout T has elapsed. This construct can be chained by several parties (up to 20 in the modern Lightning Network (cite)) to create a multi-hop payment within the network. One implication of this security model is that each party must ensure that their outgoing hash lock puzzle’s absolute timelock T_o is offset from the incoming absolute timelock T_i by a value of C_{delta} . This value C_{delta} is commonly referred to as the CLTV delta (cite bolt). This value C_{delta} is an important security parameter, as if C_{delta} blocks passes and the outgoing hash lock isn’t fully resolved, then a *race condition* occurs as the time out clauses of *both* the incoming and outgoing hash locks have expired.

Routing Nodes as Profit Seeking Capital Allocators

Entities on the Lightning Network that exist primarily to collect fees for successfully forwarding payments are referred to as *routing nodes*. A routing node commits capital to the network within payment channels in order to be able to facilitate payments in the network. (move much of that intro about the nodes here inseed?). As routing nodes incur an opportunity cost by committing capital to the network, they request a fee F upon completion of a successful payment forward. This fee $F = F_{base} + F_{rate} * A_{sat}$ is comprised of two parts: a proportional amount (a rate) and a fixed amount, which are both expressed in *milli satoshis* which are 1/1000 of the base satoshi unit.

Note that routing nodes are not compensated on an ongoing basis, and are not compensated for anything other than a completed payment. As a result, many routing nodes (cite aviv stuff) may be allocating capital in a non-productive manner as they’ve speculatively opened channels to areas of the network where no true transaction demand exists. If the Lightning Network was a physical transportation network, then it would be as if eager contractors started building roads to seemingly random destinations, only to find that those roads weren’t actually demanded at all. This information asymmetry (where new channels are actually demanded) and the current inability for today’s network participants to exchange these key demand signals lies as the crux of the bootstrapping problems of the Lightning Network.

3.2 Bootstrapping Problems in the Lightning Network

In this section, building on the background provided above, we aim to detail the various bootstrapping problems that exist in the Lightning Network today. These problems will serve motivation for our solution, the Channel Lease Marketplace, and a specific instantiation of such a construct: Lightning Pool.

3.3 New Routing Node Bootstrapping

As the Lightning Network is a fully collateralized network, in order to *join* the system, a participant must commit capital in the form of Bitcoin charged into payment channels on the network. Routing nodes, however, are in a unique situation, as they need to both *commit* their own capital to the network, as well

as *solicit* committed capital from *other* routing nodes. This is due to the fact that in order to be able to forward a payment of size P_{sat} , the routing node must first have $P_{sat_{out}}$ satoshis committed as *outbound* payment bandwidth (to use for sending) and $P_{sat_{in}}$ committed as *inbound* payment bandwidth, with the difference of the two amounts, $F = P_{sat_{out}} - P_{sat_{in}}$ being collected as a forwarding fee upon payment completion. This *pair-wise* capital commitment requirement is commonly cited as a major barrier to Lightning Network adoption (cite), as well as why large "hubs" are inherently economically inefficient (cite joseph SB montreal).

A routing node operator faces two key questions when attempting to join the network in a productive manner, while also attempting to optimize for *capital* efficiency:

1. *Where* should I open channels (thereby committing outbound capital) within the network in order to *maximize* the velocity of transactions through my channels, along with the corresponding fee revenue F_r ?
2. *How* can I attract *other* routing node operators to commit capital to my node such that I can actually forward payments to earn *any* revenue F_r ?

We argue that the above two questions, optimizing for capital efficiency and velocity of committed channels, can only properly be addressed by the *existence* of a *marketplace* that allows agents (routing node operators) to communicate their preferences using demand signals. Intuitively, a channel open to an undesirable location (possibly over-served) will have low transaction velocity C_v , and result in overall lower total fee revenue F_r . In order to maximize both C_v and F_r , a routing node should only open channels to where they're *most demanded*. If an agent is willing to pay up to $P_{premium}$ Bitcoin for inbound bandwidth, then they must gain more utility than the paid premium $P_{premium}$, as otherwise, such a transaction would not be economically rational. Thus, the existence of a marketplace that allows routing nodes to efficiently commit their outbound capital, as well as *purchase* new inbound capital is a key component to solving the bootstrapping problem for routing nodes.

3.4 New Service Bootstrapping

If routing nodes are the backbone or highway of the Lightning Network, then so called Lightning Services, are the primary *destinations* for a given payment. For simplicity, we assume that a given Lightning Service is primarily a payment *sink*, in that it's primarily *receiving* over the LN. Eventually, it may become common for a service to be balanced in terms of sending *and* receiving, resulting in a net-flow of zero, but today in the network, most flows are uni-directional (cite), creating the need for on/off chain bridges such as Lightning Loop (cite).

Demand for Incoming Bandwidth

Focusing on the case of a Lightning Service that's primarily a *payment sink*, in order to receive up to N Bitcoin, the service requires S_b Bitcoin to be committed as inbound capital, with $S_b > N$. Otherwise, assuming only channel churn (cite bryan blog post), all inbound bandwidth will become saturated, rendering a service unable to receive additional Bitcoin over the LN. Therefore,

the operative question a service operator needs to ask when attempting to join the network is:

- How can I solicit enough inbound bandwidth to be able to receive up to S_b Bitcoin?

Preference for Quality of Bandwidth

It's important to note, that as operating a *valid* routing node on the network requires a degree of skill and commitment (cite bryan blog post), some routing nodes are able to provide more *effective* service than others. As an example, imagine a routing node *Bob*, who has sufficient capital committed to his node in both the inbound and outbound directions, but who is chronically *offline*. As a node must be *online* in order to be able to forward payments, any capital C_{Bob} committed by *Bob*, can essentially be considered dead weight. With this insight in mind, we revisit the bootstrapping questions of the Lightning Service to also require a high *quality of service*:

- How can I solicit enough *high quality* inbound bandwidth within the network to be able to receive up to S_b Bitcoin?

Time Committed Incoming Bandwidth

However, from the point of view of an active Lightning Service, just having sufficient high quality inbound bandwidth may not be enough. Consider that a high quality node *Carol* may erroneously decide to commit capital elsewhere, resulting in overall lower channel velocity C_v for their channels. This type of fair-weather behavior serves as a detriment to our Lightning Services; they're unable to properly plan for the future, as they don't know *how long* the inbound bandwidth will be available for receiving payments. As a result, it's critical that the Lightning Service has a hard guarantee with respect to *how long* capital will be committed to their node. Taking this new criteria into account, we further revisit our new service bootstrapping problem statement:

- How can I solicit enough *high quality* inbound bandwidth to be able to receive up to S_b Bitcoin, that will be committed for *at least* time T_{blocks} ?

Summarizing, in addition to the existence of a marketplace for buying and selling capital commitment obligations, a would-be buyer requires some sort of *rating-system* to reduce information asymmetry (distinguish the good nodes from the lemons), and also requires that any capital committed must be committed for a period of T_{blocks} . These new requirements argue for the existence of a Node Rating agency, as well as somewhere to ensure capital will be committed for a set period of time in a trust-minimized manner.

3.5 End User Bootstrapping

Finally, we turn to the end users of the system. In our model, the end users of the system are those that are frequently *transacting*. If routing nodes are the highways in our payment transportation network, with Lightning services as popular destinations, then users trigger *payment flows* that traverse the backbone created by routing nodes, to arrive at the Lightning services. Note that

within our model, we permit end users to both send *and* receive. Compared to bootstrapping a new user to a Layer 1 system such as the Bitcoin blockchain, bootstrapping to a Layer 2 system like the Lightning Network presents additional challenges. The core challenge is created by the constraint that in order for a user to send K_s Bitcoin, they also need K Bitcoin committed within the network. Similarly, in order to receive up to K_r Bitcoin, they need up to K_r Bitcoin committed as inbound bandwidth.

From the perspective of attempting to achieve a similar user-experience as a base Layer 1 system, the receiving constraint is the most challenging. Notice that a user cannot simply download a Lightning wallet and start receiving funds. Instead, they need to first solicit *inbound* capital to their node first. Many wallet providers such as Breez and Phoenix have started to overcome this issue by committing capital to the users themselves (cite). This is essentially a customer acquisition cost: by providing this inbound bandwidth to users, the wallet becomes more attractive as it enables both sending and receiving. However, just receiving isn't enough. A user needs to be able to send *and* receive. In addition to this required symmetry, a typical user also has all the same quality of service, and time-committed capital requirements as well.

With this background, we can phrase the end user bootstrapping problem as follows:

- How can a new user join the Lightning Network in a manner that allows them to both send and receive to relevant destinations in the network?

3.6 Market Design & Auction Theory

In this section, we make a brief detour to the economic field of auction design to examine how similar resource allocation problems have been addressed by market design in existing industries. These examples include both digital and physical goods. In the modern age, market design and proper construction of corresponding *auctions* can be used to improve resource utilization and capital efficiency (cite cramton <http://www.cramton.umd.edu/papers2010-2014/cramton-market-design.pdf>). Within a particular domain, context-specific design decisions can be made so as to better optimize resource allocations for all participants. Common uses of auction design in the wild include wireless spectrum auctions by the Federal Communications Commission (FCC) (cite), package auctions for auctioning off takeoff and landing rights at airports (cite), real-time electricity markets (cite), and also carbon credits (cite). Market design bridges both theory and practice in order to solve real-world resource allocation constraints (cite <http://www.cramton.umd.edu/papers2010-2014/cramton-market-design.pdf>).

A commonly used tool in the field of market design is the concept of *auctions*. Auctions allow agents to gather and exchange pricing signals in order to determine who gets which goods, and at which price. The design of a proper auction for a particular resource allocation problem has a vast design space. For example, should a first or second price auction be used (cite)? How frequently should the auction run? What *type* of auction should be used? Should participants be able to see the bids of other participants? And so on.

Building off the series of bootstrapping problems posed above, we turn to market design as a tool to efficiently allocate our scarce resource in question:

inbound payment bandwidth. Our problem-space is unique, however, in that as we’re dealing with the allocation of capital, there are inherent opportunity costs: why should a routing node commit capital to the Lightning Network, compared to some other asset that has a similar risk adjusted rate of return (cite)? In this context, our end solution may take the form of a *money market*, which is used by entities to trade short-term debt instruments.

3.7 Money Markets & Capital Leases

In traditional financial markets, money markets are used to allow entities to trade short term debt instruments. Examples of such instruments include U.S Treasury Bills (cite), certificates of deposit (site), and repurchase agreements (repos)(cite). Capital markets on the other hand, are the long-term analogue of money markets, in that they deal with longer timeframes, and also are more heavily traded on secondary markets with retail traders being more involved.

In the context of the Lightning Network, our concept of capital obligations appears similar to a bond, in that we require a period of time for which capital is allocated. However, unlike a bond, the committed funds can only be used on the Lightning Network to provide a new type of service: the *propensity* to receive or send funds on the network. As a result, we don’t require funds in channels to be borrowed, instead they only need to be *leased* for a period of time. Also unlike bonds, wherein it’s possible for the issuer of a bond to *default*, thereby failing to repay the borrowed money, in the context of the LN, there is *no inherent default risk*. Instead, arguably the concept of channel leases can be viewed as a risk free rate of return (cite) in the context of Bitcoin, and specifically in the context of the Lightning Network.

The existence of a *channel lease* serves to provide routing nodes with an additional monetary incentive (in the form of a premium paid by the lessee of the coins) to operate a routing node. As a result, we can modify the revenue R_c earned by a routing node within the context of a specific channel C , for a period of t blocks to be:

$$R_c(t) = (P_c * t) + (F_c * P_{f_c}) \quad (1)$$

where P_c is the current *per-block* interest rate, and $(F_c * P_{f_c})$ is the *expected* routing fee revenue of the channel within that interval

We reference an *expectation* for fee revenue, as fees are effectively a *speculative* component of the routing revenue of a node. If a channel was allocated to a node in high demand, one would expect the latter portion of the question to possibly dominate the premium. If the opposite is the case, then a routing node would derive most of its revenue from the yield earned by leasing a channel. In this manner, the existence of a concept such as a channel lease actually serves to *reduce* the variance in a routing node’s revenue, similar to how joining a mining pool can reduce the variance of a Bitcoin miner’s earnings (cite).

Finally, we argue that the existence of a channel lease that pays a premium based on a *per-block* interest rate would result in a novel low-risk yield-generating instrument for the greater Bitcoin network. Such a per-block interest rate r_{b_i} would serve to allow market participants to effectively price the cost of capital on the Lightning Network. Assuming the existence of varying durations

D_1, \dots, D_n , a *yield* curve conveying the relative short and long term interest rates of channel yields could be constructed. Such an instrument would then potentially serve as the basis for higher level structured products and derivatives built on top of the base channel lease instrument.

4 Bootstrapping Problems as Solved by CLM

Prior to outlining our design for channel liquidity markets, we seek to provide a set of illuminating use cases in order to demonstrate the need for such markets on the Lightning Network. Each of these use cases are empirical real-world problems related to the Lightning Network, and prior to the publication of this document, no known solution has been presented.

4.1 Bootstrapping New Users via Sidecar Channels

A common question posed concerning the Lightning Network goes something like: Alice is new to Bitcoin entirely, how can she join the Lightning Network without her, herself, making any new *on-chain* Bitcoin transactions? It's desirable to present a solution to such a use-case as on-boarding for an on-chain user is as simple as sending coins to a fresh address. In order for off-chain payment channel networks to achieve wide-spread usage, a similar, seamless on-boarding flow must be exist.

We frame the solution to this use-case in our model of channel liquidity markets. In this case, Alice is a *new* user to the network that requires *inbound* and *outbound* liquidity. Without outbound liquidity, she's unable to send to any other node on the network. Without inbound liquidity, she's unable to receive any further payments on the network. "Sidecar channels" allow an acquaintance of Alice, let's call her Carol, to engage in a protocol with an existing routing node on the network, Bob, to provide *both* inbound and outbound liquidity for Alice. Carol is able to provide liquidity either with an off-chain, or on-chain payment. At the end of the engagement, Carol has provided channel liquidity to Alice via Bob, who himself is compensated accordingly for his role in the protocol.

4.2 Demand Fueled Routing Node Channel Selection

A "routing" node on the Lightning Network is a node categorized as having a persistent publicly reachable Internet address, a set of inbound channels from leaf nodes, and one which seeks to actively facilitate payment flows on the network in order to gain fee revenue using existing liquidity as bandwidth for these payment flows. A common question asked in the initial bootstrapping phase of the Lightning Network by node operators is: "where should I open my channels to, such that they'll actually be routed through"? We posit that channel liquidity markets are the answer to this question.

The channel liquidity market answer to this question is supplementary to autopilot [?] techniques for automated channel creation based on static and dynamic graph signals. A key drawback of autopilot channel establishment techniques is that for the most part, they're devoid of *economic* context. A particular location in the sub-graph may be "fit" or attractive from a graph

theoretic perspective, but may not lead to a high velocity channel, as there was no inherent demand for a channel created at that particular location. Using a CLM, a node operator can enter a targeted venue to determine what the *time value* of his liquidity is on the network. New services such as exchanges or merchants on the network can bid for the node's liquidity in order to serve their prospective customers, with the node earning a small interest rate up-front for committing his liquidity in the first place (scaled by the worst-case CSV delay).

4.3 Bootstrapping New Services to Lightning

Any new service that wishes join the Lightning Network faces the same problem: "How can I incentivize nodes to create *inbound* channels to me in order to be able to accept payments?". CLMs provide an elegant solution. The merchant/exchange/service uses their existing on-chain funds to enter the liquidity marketplace in order to exchange their on-chain coins for off-chain coins. Once the trade has been atomically executed, the merchant immediately has usable inbound liquidity that can be used to accept payments from users.

The merchant can then use its prior marketplace exchange partners as "introduction" points. As it's undesirable for all users to connect directly to the merchant for payment connectivity, users can instead establish channels to the prior swapping partners of the service. As the merchant acquires more liquidity in the future, they further contribute to the path diversity and strength of the network, allowing users to pay via several introduction points using AMP [?].

4.4 Cross-Chain Market Maker Liquidity Sourcing

As traders are becoming more aware of the counterparty risk of trading on centralized exchanges, they are looking to non-custodial exchange venues. The flexibility of channels on the Lightning Network appears to be a prime candidate of such a venue: channels allow for non-custodial trading at similar execution speeds to that of centralized exchanges. Additionally, channel based non-custodial exchanges are not vulnerable to front-running tactics executed by miners. Instead, the trade execution and even the prior trade history are *only* known to the participants, providing a greater degree of financial privacy.

Once again, we encounter a bootstrapping issue. How is a market maker on a payment channel-based non-custodial exchange meant to gather an initial pool of liquidity to service orders? We see CLMs as a natural solution. The market maker can seek out liquidity for relevant trading pairs by purchasing inbound channel liquidity, in addition to putting up its own outbound channel liquidity to other market makers. A balanced distribution of liquidity amongst market makers allows for new traders to participate in the exchange, knowing that their flows are balanced, meaning they can receive as much as they can send via the market maker, allowing them to instantly start to execute cross-chain atomic swaps.

4.5 Instant Lightning Wallet User On Boarding

Wallets commonly face the UX challenge of ensuring a user can receive funds as soon as they set up a wallet. Some wallet providers have chosen to open new inbound channels to users themselves. This gives users the inbound bandwidth

they need to receive, but can come at a high capital cost to the wallet provider as they need to commit funds at a 1:1 ratio. A CLM like Lightning Pool would allow wallet developers to lower their customer acquisition costs, as they would need to pay only a percentage of the funds to be allocated to a new user. Just like the merchant purchasing inbound liquidity in the above segment, a wallet provider could pay something like one thousand satoshis to have one million satoshis allocated to a user, instead of fronting the entire one million satoshis themselves.

4.6 Variance Reduction in Routing Node Revenue

Today, routing node operators aim to join the network in order to facilitate the transfer of payments as well as to earn fees over time by successfully facilitating payments. However, if a node isn't regularly routing payments (thereby earning a forwarding fee), then they aren't compensated for the various (though minor) risks they expose their capital to. With Pool, routing node operators are able to ensure that they're consistently compensated for their cost of capital.

5 The Channel Lease Marketplace

In this section, we present an overview of the Channel Lease Marketplace architectural design. In section 6, we make a brief detour to define the *Shadow Chain* application frame work, before presenting a concrete instantiation of a CLM, in the form of **Lightning Pool**.

5.1 High-Level Description

First, we describe our solution at a high-level. Drawing heavily from existing market auction design, we're interested specifically in double-call auctions which allow both the buyer and the seller to buy/sell *indivisible* units of the good in question, which in this case is a channel lease. We then build upon this base double-call auction by leveling the information playing field (cite) by making all orders *sealed-bid*. Rather than the auction being cleared continually within a central-limit order book, we instead opt to utilize a discrete interval frequent batched auction in order to mitigate front-running and other undesirable aspects. Rather than participants paying what they bid (commonly called a pay-as-you-bid auction (cite)(cite)), all participants will instead pay the same *uniform clearing price* (cite). Finally, all operations that result from a successful auction are batched and committed in a *single* atomic blockchain transaction.

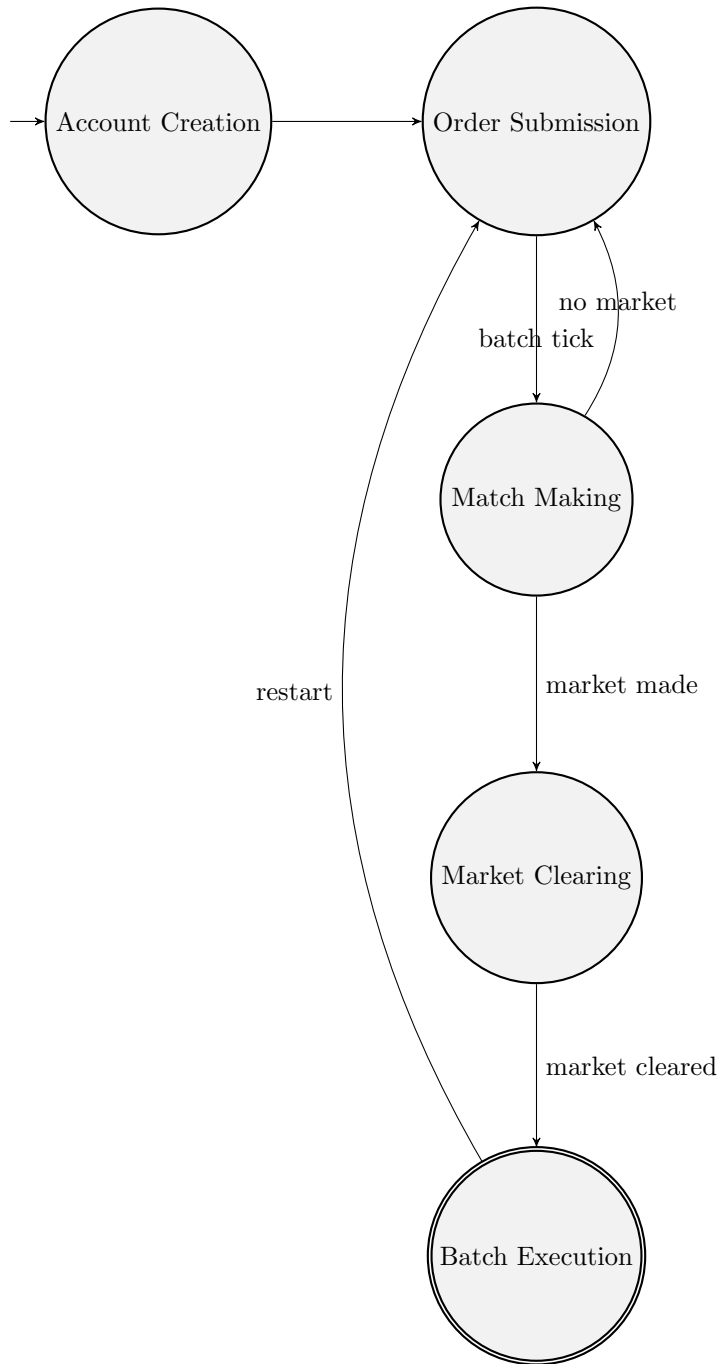


Figure 1: Auction State Machine

Marketplace Auctioneer

We assume the existence of a *non-trusted* auctioneer Λ that publishes a master auctioneer key A_{auction} ahead of time. The auction itself is uniquely iden-

tified by A_{auction} from the perspective of the system due to the *Shadow Chain* qualities of the system. The auctioneer implements a non-custodial auction via **Marketplace Accounts** which use a new unique key derived from A_{auction} as the second public key in the 2-of-2 multi-sig. The auctioneer accepts and validates orders off-chain, aides an agent in modifying their account before expiration, proposes a valid batch to each of the agents matched in an instance of the auction, and produces a batch execution transaction which modifies accounts appropriately, and creates a series of corresponding channel leases.

Account Creation

Before being able to participate in the marketplace, we require that an agent first create a **Marketplace Account**. A **Marketplace Account** is a non-custodial account that forces an agent to commit capital in the form of Bitcoin to the market for a period of time. As we require agents to fully back all orders within an account, we eliminate a number of order spoofing vectors. Additionally, the time-locked non-custodial nature of the account ensures a user is able to recover their funds fully without any additional on-chain transactions (aside from the sweeping transaction).

Marketplace Order Units

We abstract over the base satoshi unit and define a **unit** from the PoV of the marketplace which is the base unit in which all orders are expressed and settled in. We assume that the value of a given **unit** is set such that even a single lease of the smallest unit is still economical from the perspective of the base blockchain and on-chain fees. All orders *must* be divisible by a whole unit, and the final clearing volume of a given batch is also expressed in units.

Order Submission

Once an agent has created a valid **Marketplace Account**, they can enter the order submission phase. It's important to note that this order submission takes place *off-chain*. Only the final execution of an auction batch takes place on-chain. During the order submission phase, agents are free to modify their accounts and orders. Only valid orders will be accepted to be eligible for the next auction iteration.

Market Clearing

Every Υ minutes, the auctioneer attempts to *clear* the marketplace. An auction can be cleared if the lines of supply and demand cross such that at least a single unit is bought/sold. As the market has no explicit closing time, it's possible that during a market epoch, no market can be made. In the scenario that a market can be made, then rather than each participant pay what they bid, the auctioneer instead uses a single clearing price based on the market's clearing price algorithm.

Batch Execution

Once a market has been cleared, we enter the batch execution phase. During this phase, the auctioneer sends a batch proposal Π , which describes the proposed market clearing structure. Π may either be a plaintext description of a valid clearing solution, or a more "argument" describing one. Valid batches are then bundled into a single **Batch Execution Transaction** that updates all involved accounts, and creates any channel leases bought/sold in the batch. After a period of time Υ has elapsed, the market is restarted with any new orders and account being considered for market clearing.

5.2 Lightning Channel Leases

Liquidity Maker & Taker

We begin by introducing the concept of a **Liquidity Taker**:

Definition 5.1. (Liquidity Taker). A Liquidity Taker is an agent in a Channel Liquidity Market seeking to *obtain* new *inbound* channel liquidity of size A_{sat} for a period of T_{block} Bitcoin blocks.

A taker is prepared to either bootstrap the inbound liquidity with their own on-chain coins, or pay a *premium* in order to receive a "lease" of liquidity from another agent in the market. Takers populate the *demand* side of our market. They require new inbound liquidity in order to be able to immediately receive payments on the network, or to better position themselves as a routing node within the network.

A natural companion to the Liquidity Taker agent within a CLM is the **Liquidity Maker**:

Definition 5.2. (Liquidity Maker).

A Liquidity Taker is an agent in a Channel Liquidity Market seeking to earn *yield* by deploying up to A_{sat} Bitcoin into the Lightning Network for up to a period of T_{block} Bitcoin blocks, earning a profit α .

Notice that we utilize Bitcoin block-time rather than wall-clock time (Median Past Time (cite)) in these definitions, as we seek to enforce these durations using Bitcoin Script and using block-time is more objective compared to wall-clock time.

The profit (α) earned by a **Liquidity Maker** takes two forms:

- A one-time *premium*, $R_{premium}$, commanded by the Maker which reflects the latent demand and time-value of regular coins vs "lifted coins" (coins placed in channels).
- Ongoing recurring revenue, F_c , in the form of forwarding fees earned by facilitating payments to their matched taker.

We argue that the existence of such Channel Liquidity Markets will increase the *efficiency* of capital deployed to a payment channel network by allowing agents to signal the relative demand of lifted coins compared to non-lifted coins. Additionally, such markets also allow an existing routing node on the network to *re-allocate* lifted coins from a low-velocity section of the sub-graph, to one of higher velocity:

Theorem 5.1 (Channel velocity revenue). Holding all channel liquidity equal, channels allocated to a higher velocity section of the sub-graph will yield a higher F_c than channel allocated to a low-velocity section of the sub-graph.

Intuitively, if each payment flow sourced at an incoming channel C_i and sunk at an outgoing channel C_o pays an equal forwarding fee per flow, then for a fixed unit of time, a higher velocity channel will result in higher total revenue in a time slice.

The role of Channel Liquidity Markets in a payment channel network is to reduce information asymmetry by allowing agents to signal their preferences for lifted coins vs non-lifted coins. The existence of *venues* where these markets can be carried out benefits the wider network by allowing agents to determine where their liquidity is most demanded on the network.

Channel Leases

With our two primary agents defined we now move on to the definition of a Lightning Channel Lease:

Definition 5.3. (Lightning Channel Lease). A Lightning Channel Lease is defined as, $\Gamma = \{P_T, P_M, A_{sat}, D_{block}, r_i\}$, where:

- P_T is the `secp256k1` public key (cite) of the **Liquidity Taker**, P_M the public key for the **Liquidity Maker**.
- A_{sat} is the total amount of Bitcoin within the contract.
- D_{block} is the duration of the contract expressed in Blocks.
- r_i is the per-block interest rate as discovered in the i th instance of the market.

Note that the premium R_P as referenced above is parametrized in using the *lease duration* D_{blocks} : $R_P(D_{blocks}) = r_i * D_{blocks}$ as we deal in simple, rather than compounding, interest. The duration of the contract D_{blocks} is of great interest as similar to U.S Treasury auctions, a *yield-curve* (cite) can be constructed based on the matched contents of a given auction iteration.

5.3 Non-Custodial Auction Accounts

In order to participate in the auction, we require all participants to deposit their trading balance into a Marketplace Account:

Definition 5.4. (Marketplace Account). A marketplace account is a *non-custodial* account defined as, $\Psi = \{K_{sat}, T_{blocks}, P_{acct}, \Omega_{nodes}\}$ where.

- K_{sat} is the total amount of Bitcoin available within the account.
- T_{blocks} is the absolute expiry height of the account.
- P_{acct} is a `secp256k1` public key that *uniquely* identifies the account.
- Ω_{nodes} is a set of Lightning Network nodes controlled by the account.

We stress that these accounts are *non-custodial* in that after a period of time T_{blocks} the agent is able to freely remove the funds from their account. Before this period has passed, an agent may require the aide of the auctioneer to close, deposit, or withdraw funds from their account. In the case of the **Liquidity Taker**, the funds within the account K_{sat} , must be enough to pay for any desired premium. Conversely for the **Liquidity Maker**, we require all funds they wish to lease out to be deposited into the account.

This structure, which forces all participants to fully commit all funds they wish to use within the marketplace into a non-custodial account, is similar to the concept of Fidelity Bonds (cite). This structure has a number of desirable properties including:

- **Order spoofing mitigation:** Within the CLM, as all orders must be "fully backed", it isn't possible to place a "fake" order that cannot be filled.
- **Time value opportunity cost:** By forcing all agents to suspend funds they wish to use within the market, those funds cannot be used elsewhere, thereby adding an implicit cost to joining the marketplace.
- **Deterministic batch execution construction:** As we'll see in later sections, the existence of a fixed account for each agent simplified the clearing and execution process within the auction lifecycle.

These accounts in the abstract may take many forms, but as we focus on Bitcoin, as detailed in later sections, these accounts will take the form of a multi-sig output, with one key belonging to the auctioneer.

5.4 Order Structure & Verification

With our channel lease contract and account structure defined, we now move on to our order structure. As with any auction, orders are how the agents express their *preferences* with respect to what they wish to buy and sell. Importantly, all orders within the market must be backed by a valid non-expired account, and must carry an *authentication tag* which prevents order spoofing, and also ensure proper integrity of a given order once it has been submitted.

Order Structure

We define an **Order** within the context of a CLM as follows:

Definition 5.5. (Order). An **Order** is a authenticated n -tuple:

$\Theta = \{O_{type}, K_{nonce}, V_{ver}, P_{acct}, \Delta_{base}, \Delta_{aux}, T_{auth}\}$, where:

- $O_{type} \in \{\text{Ask}, \text{Bid}\}$ denotes if an order is an Ask or a Bid. In addition to the version, this may affect how the Δ_{aux} attribute is parsed.
- K_{nonce} is an *order nonce* which uniquely identifies this order, and is typically derived as $K_{nonce} = \leftarrow \$\mathbb{Z}_p$.
- V_{ver} is the *version* of this order. As we'll see below, the version used as an upgrade mechanism, and is needed in order to parse any newly added fields, as well as compute the digest required to check the authentication tag attached to an order.

- P_{acct} is the public key that uniquely identifies this account.
- The set of base order details is:
 $\Delta_{base} = \{\alpha_{rate}, A_{sat}, M_{pub}, L_{pub}, A_{addr}, C_{type}, D_{blocks}, F_{chain}\}$, where:
 - α_{rate} is the desired *per-block* rate that the owner of the order wishes to buy/sell a channel lease at. Further below, this may be referred to as the BPY or block percentage yield.
 - A_{sat} is the total contract size expressed in lease *units*. Restricting orders to whole units simplifies preference matching within the system.
 - M_{pub} is the multi-sig public key to be used when creating the funding output (cite) of the transaction.
 - L_{pub} is the identity public key (cite) of the Lightning Node that wishes to buy/sell this channel.
 - A_{addr} is the network address to be used to connect to the backing L_{pub} to initiate the channel funding process if this order is matched.
 - C_{type} is the *type* of channel to be created if this order is matched.
 - D_{blocks} is the target *lease duration* of the contract.
 - $F_{chain_{max}}$ is the max chain fee expressed in *sat/byte* that the owner of said order is willing to pay within a batch.
- The set of *auxillary* details is implicitly defined by the order version V_{ver} .
- T_{auth} is an authentication tag that allows the auctioneer, and other traders to validate the integrity and authenticity of the order.

An order allows a **Liquidity Taker** or a **Liquidity Maker** to express their *preference* with respect to what type of channel lease they’re looking to buy/sell.

Order Validation

Returning back to our tag T_{auth} , we will now specify how such a tag is to be computed, and verified. In the abstract, we require that the tagging scheme is SUF-CMA secure (cite). Given this security requirement, we define two polynomial-time algorithms: (**GenOrderTag**, **VerifyOrderTag**) with the following requirements:

- **GenOrderTag**($P_{acct_{priv}}, \Theta$) $\rightarrow T_{auth}$. Given an input of the private key that corresponds to the public key of an account, and the complete order details, a valid tag T_{auth} is generated.
- **VerifyOrderTag**($P_{acct}, \Theta, T_{auth}$) $\rightarrow b$. Given a public key of an account holder, a valid tag, and the order itself, **VerifyOrderTag** outputs $b = 1$ if the tag is valid.

As we use a public-key based tagging technique, the validity of an order is verifiable by any other active trader within the marketplace including the auctioneer of the market place.

5.5 Auction Design

In this section, we describe the abstract definition of a **Channel Liquidity Marketplace**, which addresses each of the issues presented in the bootstrapping section of the background, by creating a new form of batched auction which allows **Liquidity Takers** and **Liquidity Makers** to buy/sell Lightning Channel Leases in a non-custodial manner.

5.5.1 Auction Specification

In this section, we'll now specify the behavior and requirements of an using abstract **Channel Liquidity Marketplace** instance. We define the expected behavior and the client-facing interface of a CLM instance. A CLM is a tuple of polynomial-time algorithms divided into five distinct but related categories:

- System Initialization: **SystemInit**
- Account Operations: (**NewAccount**, **ModifyAccount**)
- Order Book Maintenance: (**SubmitOrder**, **CancelOrder**)
- Market Clearing: (**MatchMake**, **MarketClearingPrice**, **ClearMarket**)
- Batch Execution: (**ConstructBatch**, **ExecuteBatch**)

With behavior and semantics as expressed below.

System Initialization

Before the market place can be used, we require it to be initialized by the auctioneer. This initialization is a one-time process, and doesn't result in any trapdoor or "toxic waste" material being produced:

SystemInit($1^\lambda, \Upsilon_{min}$) $\rightarrow (P_{auction_p}, P_{auction_s}, \Psi_A)$. Denoting the security parameter as λ , the **SystemInit** algorithm takes as input the security parameter, and the batch interval Υ_{min} expressed in minutes, and outputs a public ($P_{auction_p}$) and private ($P_{auction_s}$) key pair for the auctioneer. The auctioneer's public key will be used as an parameter in algorithms related to account creation, modification, and batch execution. This algorithm also returns Ψ_A , which is a special account owned by the auctioneer that will be used to collect fees, and during batch construction.

Account Operations

In order to create an account, agents will need to interact with the auctioneer itself. After account creation an account can freely be modified (close, deposit, withdraw, etc) if the account isn't part of an active batch:

NewAccount($1^\lambda, P_{auction_p}$) $\rightarrow \Psi$. The **NewAccount** algorithm takes as input our security parameter, and the auctioneer's public key, and outputs a new account for the new agent within the marketplace. We require that all resulting accounts within the marketplace be *unique*. We permit a single logical agent to

have multiple accounts.

$\text{ModifyAccount}(\Psi, P_{\text{auction}_p}) \rightarrow (b, \Psi')$. The **ModifyAccount** algorithm takes an existing valid account Ψ and the auctioneer's public key and performs an account modification. The algorithm returns $b = 1$ if the modification was successfully, and $b = 0$ otherwise. An account modification may fail if the target account is already part of a pending batch. An account modification can either:

- Deposit new coins into the account.
- Withdraw coins from the account.
- Close the account by removing all coins from the account.

Note that as each of these operations require an on-chain transaction, they can freely be batched with other on-chain transactions, or even the transaction that executes an auction's batch.

Order Book Maintenance

Once accounts in the marketplace are open, agents are able to submit orders between batch epochs. The *size* of all orders is expressed in units, and as we mention below, we permit partial matches of an order. A partial match can either update the order state in place, or require the agent to re-submit a new valid tag for the modified order in the batch execution phase:

$\text{SubmitOrder}(\Theta, T_{\text{auth}}) \rightarrow b$. The **SubmitOrder** algorithm takes as input a structurally sound order Θ , and its authentication tag T_{auth} and outputs a bit b . The bit $b = 1$ if the order is valid according to market place rules, and the **VerifyOrderTag** returns $b = 1$ given the specified parameters.

$\text{CancelOrder}(\Theta, K'_{\text{nonce}}) \rightarrow b$. The **CancelOrder** given an existing order Θ and the *opening* of the K_{nonce} commitment K'_{nonce} and returns $b = 1$ if the commitment opening is valid, and there exists an order identified by the base K_{nonce} value.

Market Clearing

Once all orders have been placed, and the batch interval of Υ has elapsed, the auctioneer will attempt to clear the market using the following algorithms:

$\text{MatchMake}(\{\Theta_0, \dots, \Theta_n\}) \rightarrow \Phi_b = \{(\Theta_{b_0}, \Theta_{a_0}), \dots, (\Theta_{b_n}, \Theta_{a_n})\}$. The **MatchMake** algorithm takes as input the set of valid orders submitted during the past batch interval and outputs a series of tuples which reflect properly matched orders. Θ_a represents an order whose $O_{\text{type}} = \text{Ask}$, while Θ_b represents an order whose $O_{\text{type}} = \text{Bid}$. Note that since we allow *partial* matches, a given order may appear multiple times in the final match set. We require that a valid implementation be able to perform proper *multi-attribute* (cite) matching due to the existence of the Δ_{aux} portion of an order's structure.

$\text{MarketClearingPrice}(\Phi_b) \rightarrow c_{\text{price}}$. The **MarketClearingPrice** algorithm accepts the set of orders matched by the **MatchMake** algorithm and returns the

market clearing price of the prior batch. The precise market clearing price algorithm is left as a free parameter, with algorithms such as first-rejected-bid (cite) or last-accepted-bid (cite) likely being used. Utilizing of a single market clearing price is intended to promote fairness (cite) (all traders pay the same price!).

$\text{ClearMarket}(\Psi_A, \Phi_b, \{\Psi_0, \dots, \Psi_n\}, c_{price}) \rightarrow (\Psi'_A, \{\Gamma_0, \dots, \Gamma_n\}, \{\Psi'_0, \dots, \Psi'_n\})$. The **ClearMarket** algorithm takes as input a prior set of matched orders within a batch, the auctioneer's account, the set of accounts involved in the batch, and the market clearing price of a given batch and outputs: a set of channel leases to be created by a batch and a set of updated accounts which represents the state of the involved accounts post batch as well as an updated version of the auctioneer's account which may have accrued any trading fees during market clearing.

- As shorthand, we use Δ_i to refer to a cleared batch (the set of resulting accounts after the updates have been made to produce the set of desired channel leases).

Batch Execution

Once we've been able to make a market, and have the description of the resulting market state (the accounts, and the channel leases to be created), we can now move on to *executing* the resulting batch. We use the following algorithms to do so:

$\text{ConstructBatch}(\Delta_i) \rightarrow B_{t_i}$. The **ConstructBatch** algorithm takes a valid market clearing (which can be seen as a delta on the auction state) and returns a valid transaction, which *atomically* executes the given batch on the blockchain.

$\text{ExecuteBatch}(B_{t_i}) \rightarrow b =$. The **ExecuteBatch** algorithm takes a fully valid batch and attempts to commit it, by confirming the transaction in the target base blockchain. Once the batch has been confirmed, all operations contained within a batch are considered executed, and can be used as inputs to additional iterations of the auction life cycle.

6 The Shadowchain: A Bitcoin Overlay Application Framework

In this section, we present the concept of a **Shadowchain**, a non-custodial application overlay framework that we'll use to construct a concrete instantiation of a CLM. We note that shadowchains may also be of independent interest, as they're a novel way to layer more complex interactions on top of the base Bitcoin blockchain. We note that shadowchains as we present them can be implemented on the base Bitcoin blockchain today without any additional changes or enhancements. However, further extensions to Bitcoin such as non-interactive signature aggregation (cite) and covenants (cite) can serve to dramatically increase the scalability traits of a shadowchain.

6.1 High-Level Description

First, we provide a high-level description of the shadowchain application framework.

The Shadowchain Usecase. A shadowchain can be used to implement non-custodial smart contract systems on top of the base Bitcoin blockchain. Typically one would opt for a shadowchain if the complexity of the state transition logic of the smart counteract system cannot be fully expressed using the base Bitcoin Script. Shadowchains allow an application designer to use the Bitcoin blockchain for what it's best for: censorship resistant settlement, pushing the more complex portions of the application (state, logic, etc) *off-chain*.

Shadowchain Roles & Lifted UTXOs. A shadowchain has two primary classes of agents: users, and the orchestrator. The orchestrator defines the state transition function of the shadowchain, a set of non-trusted initialization parameters, and upgrade mechanisms. A user is able to join a shadowchain by "lifting" their UTXOs *into* the higher-level shadowchain. The process of lifting (defined further below) entails the user placing funds within a time-lock released multi-sig output between itself and the shadowchain orchestrator.

Shadowchain Operation. The shadowchain orchestrator accepts transaction data from users, then periodically proposes a new *shadowchain block*. A shadowchain block takes as input the set of **Lifted UTXOs** which accepted the latest block proposals, and produces a set of *new* UTXOs, which are the end state after the state transition function has been evaluated. A shadowchain is even permitted to use *multiple* distinct state transition functions. As the funds of an end user cannot move without both multi-sig signatures, users are able to fully validate (possibly using techniques such as zero knowledge proofs (cite shafi)) that the resulting UTXO state was properly derived from the known state transition function. Note that due to this structure complex "exit-games", or fraud proofs are not required as a user simply won't sign off on a fraudulent state, and a user's UTXO is always manifested (in a base form) on the main blockchain.

Ephemeral Lifted UTXOs. In the scenario that the shadowchain orchestrator disappears, or is unresponsive, users are able convert their lifted UTXO into a regular one, by spending their coins after the time-lock has expired. This construct of an *ephemeral* lifted UTXO has a number of desirable properties on the application level, as the time-locked commitment of funds can serve to mitigate a number of application-level issues such as spam or sybil resistance.

Shadowchain Cut-Through As the evolution of a state transition function happens *off-chain*, it's possible to coalesce several distinct shadowchain blocks into a single block which is the composition of the successive invocations of the state transition function. This technique is similar to transaction cut-through (cite), but performed in a multi-party setting. Leveraging this technique, the shadowchain orchestrator can optimistically treat the current latest shadowchain transaction in the mempool as an in-memory data structure to be updated off-chain (via transaction replacement techniques (cite)), with the state being "written to disk" once confirmed. As a result, it's possible to commit several shadowchain states (possibly hundreds) in a single logical Bitcoin transaction.

Shadowchain Upgrades. Finally, similar to the base blockchain, a shadowchain can also be *upgrade* in a forwards and backwards compatible manner. In other words, it's possible for a shadowchain orchestrator to *soft-fork* the state

transition logic by restricting a valid state transition to enable new behavior. Notably, the operator can do this in a de-synchronized manner as only those wishing to use features in the new state transition function need to adhere to the new rules. Additionally, an operator can opt to also introduces new backwards incompatible state transition functions. Due to the inherent batched nature of Bitcoin transaction, an operator can commit multiple logical shadowchain blocks (with distinct state transition functions) in a single atomic Bitcoin transaction.

To summarize, the shadowchain application framework is a novel technique for constructing overlay applications on the base Bitcoin chain in a non-custodial manner. Shadowchains avoid the complexity of fraud proofs and exit-games by ensuring that the user has custody of their funds at all times and is able to fully validate any proposed state transition. Shadowchains are able to compress several logical state transitions into a single Bitcoin transaction using a multi-party cut-through technique. An orchestrator of a shadowchain is also able to upgrade the state transition logic on the fly, in a purely off-chain manner.

6.2 Comparison To Related Frameworks

6.3 The Shadowchain Framework

In this section, we present the abstract shadowchain application framework. Applications are intended to use this framework, providing implementations of specified virtual functions to fully specify and execute their application.

6.3.1 Shadowchain Orchestrator

First, we introduce the glue that keeps a shadowchain together, the orchestrator:

Definition 6.1. (Orchestrator). The **Orchestrator** is a non-trusted entity at the root of a shadowchain, parametrized by its long-term public key: $O_{chain} = P_O$. The duty of an **Orchestrator** is to propose new blocks (the result of a state transition) to the set of live **Lifted UTXOs** that make up the shadowchain.

A given **Orchestrator** is a non-trusted entity, and can be uniquely identified by its longer-term public key. The long-term public key P_O can also be used to uniquely identify a given shadowchain, similar to the Genesis Block hash of a normal blockchain.

6.3.2 Lifted UTXOs

Next, we define the **Lifted UTXO** which is the representation of a user's state within a given shadow chain:

Definition 6.2. (Lifted UTXO). A **Lifted UTXO** is a tuple, $\phi_U = (A_{sat}, T_{expiry}, P_u, P_o)$, where:

- A_{sat} is the size of an **L0 (Lifted UTXO)** expressed in *satoshis*.
- T_{expiry} is the *absolute* expiry height of the **L0**, after-which the owner is able to unilaterally move the funds back to the "base" Bitcoin blockchain.
- P_u is the public key of the end-user, which is 1/2 of the public keys used in the public key script of the output which manifests this **L0** on the base blockchain.

- P_o is the public key of the **Orchestrator**, typically derived from its base long-term key P_O . This key will be used as the other half of the multi-sig script of the on-chain manifestation of the LO.

The construct of a **Lifted UTXO** is similar to the existing concept of a **Fidelity Bond** (cite), yet with an application specific twist. This process is akin to creation a new 'account' within a Shadowchain. The time-lock release nature of the UTXO means a user can always recover funds if the **Orchestrator** becomes unresponsive. In addition to this, a natural cost in the form of chain fees is added which increases the barrier for potentially malicious users to interact with the shadow chain.

6.3.3 The Shadowchain

In this section, we present the abstract definition of a shadowchain, building upon the definition provided above. In addition to this, we describe the typical shadowchain life cycle using the aid of some additional helper functions, which are also intended to be supplied by the core application logic itself.

Shadowchain Components

First, we define the core components of the shadowchain.

Definition 6.3. (Shadowchain). An instantiation of a **Shadowchain** is defined as a tuple: $\Sigma = (U_L, U_O, \Delta_F, E_{exe}, A_T)$, where:

- $U_L = \{\phi_i, \dots, \phi_n\}$ is the set of non-expired **Lifted UTXOs** observed by the orchestrator.
- U_O is the current UTXO orchestrator, where they may accrue application level fees.
- $\Delta_F = \{\Delta_{f_0}, \dots, \Delta_{f_n}\}$ is the set of current state transition functions.
- E_{exe} is the abstract execution environment of the shadowchain which all participants will use to verify the correctness of a proposed state transition.
- A_T is the abstract form of the structure of the higher-level application's fundamental transaction.

Shadowchain Algorithms

Given the above components, we define the operation of a shadowchain using a series of polynomial-time algorithms segmented into the following logical categories:

- System Initialization: **InitChain**
- UTXO Management: (**LiftUTXO**, **UnliftUTXO**, **ExitChain**)
- Block Proposal & Validation: (**ConstructBlock**, **ProposeBlock**)
- Chain Execution: **CommitBlock**
- Block Cut-Through: **CoalesceBlocks**

- Chain Upgrade: **UpgradeChain**

With behavior and semantics as expressed below.

System Initialization

Before a shadowchain can be used for a given application, it is required that the system be initialized to obtain the long-term public key of the auctioneer, as well as the execution environment, and set of state transition functions:

InitChain(1^λ) $\rightarrow (U_0, P_0, \Delta_F, E_{exe})$. Given the security parameter λ (expressed in unary), the **InitChain** method returns the initial self-lifted UTXO of the orchestrator, the long-term public key of the orchestrator, and the set of initial state transition functions along with the starting execution environment.

UTXO Management

Once the chain has been initialized, given the parameters of the shadowchain, users can be lifting their UTXO to be able to participate in shadowchain blocks and operations. The process of entering the shadowchain is referred to as UTXO Lifting, while exiting is the reverse process:

LiftUTXO($T_{expiry}, \{U_{N_0}, \dots, U_{N_n}\}, P_0$) $\rightarrow \phi_U$. The **LiftUTXO** algorithm takes a series of normal UTXOs (U_N), the absolute expiration height of the UTXO, and the long-term public key of the Orchestrator P_0 and outputs a new **Lifted UTXO** for the target user. The sum of the set of input UTXOs must be greater-than-or-equal to the size of the resulting **Lifted UTXO**.

ExitChain(ϕ_U, B_{height}) $\rightarrow U_N$. Given a lifted UTXO with expiration height $T_{expiry} > B_{height}$, where B_{height} , the **ExitChain** method spends an existing lifted UTXO and resolves the user's funds in a regular unencumbered UTXO. Users will use this algorithm if they wish to exit the chain, with the Operator no longer being responsive.

UnliftUTXO(ϕ_U) $\rightarrow U_N$. Given a lifted UTXO the **UnliftUTXO** method create a new normal un-lifted UTXO output which returns all funds to the user. This is the optimistic version of the **ExitChain** algorithm, in that it requires cooperation from the orchestrator as the time-release clause of the Lifted UTXO's script is not yet unlocked.

Block Proposal & Validation

Once a shadowchain has a sufficient number of lifted UTXOs, and the system has been fully initialized, block proposal and validation can commence. This process is similar to the process of nodes broadcasting transactions, and miners ordering them within blocks in the base Bitcoin system:

ConstructBlock($\phi_{live}, T_{xn}, E_{exe}, \Delta_F$) $\rightarrow B_S$. Given inputs of the set of 'live' Lifted UTXOs, the set of transaction belonging to the live UTXOs, the execution environment, and the current set of valid state transition functions, the **ConstructBlock** outputs a valid shadowchain block to extend the current machine chain. Where:

- $\phi_{live} = live(\{\phi_{U_0}, \dots, \phi_{U_N}\})$ is the set of 'live' Lifted UTXOs where the *live* algorithm uses a heart-beat like protocol to detect the current set of active users.
- T_{xn} is the application-specific transaction format used within the shadowchain.
- $B_S = (T_{xn}, \{\phi_{U_0}, \dots, \phi_{U_N}\}, \Delta_f, \{\phi'_{U_0}, \dots, \phi'_{U_N}\}, U_A)$, is the shadowchain block itself, which is composed of the set of application transaction, input Lifted UTXOs, the resulting output UTXOs after applying the set of state transition functions, and U_A any new application-specific UTXOs produced as a result of the state transition function.

Once a block has been constructed, the operator of the shadowchain now must propose said block to the set of live Lifted UTXOs before we can move onto the next phase of shadowchain operation. As a given user may reject a block, either implicitly due to being offline, or explicitly due to a violation of the shadowchain consensus rules, then this phase may be repeated a number of times. The operator will use the following algorithm to propose blocks:

ProposeBlock(B_S, ϕ_{live}) $\rightarrow b$. The **ProposeBlock** attempts to propose the given shadowchain block to the set of live Lifted UTXOs. The algorithm returns $b = 1$ if all of the participants accept the block. Once all participants have accepted the block, we can now proceed to the execution and block commitment phase.

Chain Execution

Once the operator has established a stable set of participants which accept the proposed shadowchain block, it can execute the block and commit it in the base Bitcoin blockchain:

CommitBlock(B_S) $\rightarrow (b, TX_{id})$. The **CommitBlock** takes a valid shadowchain block, and attempts to obtain all necessary witnesses to unlock the Lifted UTXOs to re-create them post state function application along with any new application-specific UTXOs. The algorithm returns $b = 1$ if the operator was able to successfully obtain all necessary witnesses, and broadcast the Bitcoin transaction which commits the shadowchain block. Otherwise, the operator may need to re-propose a new block to a sub-set of the live Lifted UTXOs.

Block Cut-Through

Given the structure of a shadowchain block, and the state transition function itself, it's possible for a shadowchain operator to compress several distinct shadowchain blocks into a single instance that is equivalent to the successive state transition function application on each distinct shadowchain block:

CoalesceBlocks($\{B_{S_0}, \dots, B_{S_N}\}$) $\rightarrow B'_S$. Given a series of *consecutive* shadowchain blocks, the **CoalesceBlocks** algorithm compresses each consecutive block into a single block B_{S_N} which has an equivalent output state to the serial application of the state transition function on each individual shadowchain

block.

The **CoalesceBlocks** algorithm is similar to the concept of transaction cut-through (cite) for UTXO-based blockchains. Using this algorithm, the operator is able to propose a new equivalent block which produces the same set of net Lifted UTXO outputs, along with any other application specific outputs produced by any of the intermediate blocks. This can be done post-facto, and also in an optimistic manner in order to coalesce several unconfirmed shadowchain blocks into a single one.

Chain Upgrade

The process of upgrading the chain in terms of the *types* of application-level transactions offered, or the set of valid state transition functions can be done entirely off-chain in a de-synchronized manner. In order to update the environment, state transition functions, or the application-level transactions, the operator simply needs to utilize the following algorithm:

$\text{UpgradeChain}(\Delta_{new}, E'_{exe}, T'_{xn}) \rightarrow \perp$. The **UpgradeChain** is an algorithm that's executed entirely in an off-chain manner, then lets the operator of the shadowchain upgrade some or all of: the application transaction data, the execution environment, and the set of state transition functions. Due to the nature of the shadowchain, in order to utilize this new functionality each user must update their local state. However, note that the Lifted UTXO remains the same, as the operator's long-term public key hasn't changed.

6.3.4 Shadowchain Operation

In this section, we'll put together the above algorithms to outline the main execution loop of the shadowchain from the perspective of the operator as well as participants. Shadowchain operation can be viewed as a linear deterministic state machine that uses the main Bitcoin blockchain in order to transition to certain states.

Operator State Machine

The main Operator State Machine loop first attempts to process any new UTXO lifting requests, then will optimistically attempt to merge in any existing unconfirmed shadowchain blocks that can be coalesced. Independent of either of these clauses, we'll then enter into our main loop where the operator attempts to construct a new block given the set of live transaction, propose the block to the set of live UTXOs, filter out any participants that rejected the block, before finally attempting to commit the new block within the chain.

OrchestrateChain(Δ_F, E_{exe})	Oracle O
1: repeat	1: some code
2: if haveNewLiftReqs()	2: more code
3: $\{\phi_{\text{new}}\} \leftarrow \text{liftNewUTXOs}(P_O)$	
4: $\{\phi_U\} \leftarrow \phi_U \cup \phi_{\text{new}}$	
5: if numUnconfBlocks() > 1	
6: $B'_S \leftarrow \text{CoalesceBlocks}(\text{unconfBlocks}())$	
7: $(b, txid) \leftarrow \text{CommitBlock}(B'_S)$	
8: $T_{xn} \leftarrow \text{liveTransactions}()$	
9: $\phi_{\text{live}} \leftarrow \text{liveLiftedUTxos}()$	
10: $b \leftarrow 0$	
11: while $b == 0 \& \text{len}(\phi_{\text{live}}) > 0$	
12: $B_S \leftarrow \text{ConstructBlock}(\phi_{\text{live}}, T_{xn}, E_{exe}, \Delta_F)$	
13: $b \leftarrow \text{ProposeBlock}(\phi_{\text{live}}, B_S)$	
14: if $b == 0$	
15: $\phi_{\text{live}} \leftarrow \text{filterRejects}(\phi_{\text{live}})$	
16: if $b == 1$	
17: $(b', TX_{id}) \leftarrow \text{CommitBlock}(B_S)$	
18:	

Participant State Machine

The main state machine if each shadowchain participant is essentially a mirror of the orchestrator state machine. First we'll process any requests to modify our existing Lifted UTXO, gather up any unconfirmed transactions, then await a new block proposal by the orchestrator.

ExtendChain($\phi_U, \Delta_F, E_{exe}$)	Oracle O
1: repeat	1: some code
2: $T'_{xn} \leftarrow \text{unconfTxns}()$	2: more code
3: submitTxns(T_{xn})	
4: $B'_S \leftarrow \text{await } \text{recvBlockProposal}()$	
5: $b \leftarrow \text{ValidateBlock}(B'_S)$	
6: if $b == 1$	
7: sendWitnesses(U_O)	
8: $b' \leftarrow \text{await } \text{blockFinalize}$	
9: if $b' == 1$	
10: localCommitBlock()	

7 Lightning Pool: A Channel Liquidity Marketplace as a Shadow Chain

In this section, we build upon the prior sections outlining the abstract **Channel Lease Marketplace** definition, as well as shadowchain operation, and construct

out **Lightning Pool** implementation at a low-level. We first begin by detailing our implementation of the CLM algorithms, as well as our choice of certain free parameters. With this concrete structure in place, we'll then go up a layer of abstraction to demonstrate how Lightning Pool can be operated as a shadowchain on Bitcoin today, without any further modifications/enhancements.

7.1 Instantiating a CLM

7.1.1 System Initialization

Batch Key Parameter

Before an instance of Lightning Pool can be used by willing agents, the system must first be initialized. This operation can be performed only by the Orchestrator of the auctioneer. Within the system, we'll utilize incremented Elliptic Curve point which we refer to as the **batchID** for several operations. The **batchID** serves to uniquely identify a given batch, and is incremented after each successful batch.

The **batchID** itself is a nothing up my sleeve (NUMs) point which has been generated in a manner that no one, not even the auctioneer knows the discrete log to. The raw serialized **batchID** (for the very first batch) within the **Lightning Pool** system can be expressed in the following syntax (displayed using hex-encoding of a compressed key encoding of the batch key itself):

$$B_{k_0} = NUMS_{gen}(1^\lambda)$$

Let the current **batchID** for the n^{th} batch be B_{k_i} . Let G be a cyclic group of order $N = pq$, generated by an element g . We now define two helper functions to allow us to "seek" round the **batchID** key-space:

$\text{IncrementBatchKey}(B_{k_i})$	$\text{DecrementBatchKey}(B_{k_i})$
1 : return $B_{k_i} + g$	1 : return $B_{k_i} - g$
$\frac{\text{InitBatchKey}}{1 : \text{return } B_{k_0}}$	

With the system initialization, we'll now move onto specifying the structure of a **Marketplace Account** within **Lightning Pool**.

Auctioneer State Initialization

Now that we have defined the set of initialization and manipulation methods for our batch key parameter, we'll move onto the initialization of the remaining system. In order to prevent key-reuse across the entire system, we employ a similar key-derivation scheme to that of the Lightning Network's current commitment transaction format (cite). This key derivation will be used within all

account scripts within the system, as well as the auctioneer's main account.

First, we define a helper function for deriving the auctioneer's current key parameterized by the current **batchID**. As noted above, the **batchID** serves as both a public key within the system as well as a counter. The **batchID** may be expressed as a normal compressed public key, or as an integer N_b which denotes the scalar multiple off-set from the starting batch key B_{k_0} : $[B_{k_0}]N_b$. We define the **auctioneerScript** as follows:

$\text{auctioneerScript}(A_{pk}, B_{k_i})$ <hr style="width: 100%;"/> 1 : $A_{pk_i} \leftarrow A_{pk} + H(A_{pk} \ B_{k_i})$ 2 : return OP_CHECKSIG A_{pk_i}

The script itself is a simple script that simply verifies a proper auctioneer signature given a particular **batchID**. In order to make all scripts uniform (as the account scripts are P2WSH outputs, we wrap this script in a P2WSH layer of indirection as well. With this algorithm defined, we can now define the **SystemInit** implementation for **Lightning Pool** that derives the first batch key along with the starting script of the auctioneer.

$\text{SystemInit}(1^\lambda, \Upsilon_{min})$ <hr style="width: 100%;"/> 1 : $batchKey \leftarrow \text{InitBatchKey}()$ 2 : $A_{sk} \leftarrow \$\mathbb{Z}_p$ 3 : $A_{pk} \leftarrow GA_{sk}$ 4 : $aScript \leftarrow \text{auctioneerScript}(A_{pk}, batchKey)$ 5 : $pkScript \leftarrow \text{p2wsh}(aScript)$ 6 : return $(A_{sk}, A_{pk}, \text{utxo}(pkScript, seedSats))$
--

Once the third return value, auctioneer's master account has been confirmed. Participants are able to open accounts, submit orders, and finally participate in auction batches.

7.1.2 Lightning Pool Accounts

Next, we move onto the **Marketplace Account** structure for agents within the **Lightning Pool** marketplace itself. Similar to the auctioneer's master account, we apply a key derivation scheme that combines the auctioneer's key, the batch key, a key supplied by the trader, and a distinct trader specific secret to ensure, that:

- A P2WSH output script is *never* re-used.
- Traders rotate keys with each batch.
- The set of trader keys using within a batch itself is indistinguishable w.r.t the input keys referenced and newly created outputs keys.

Let P_t be a trader's base key, $S_t \leftarrow \$\mathbb{Z}_p$ be an account-specific secret generated by the trader, and P_A be the auctioneer's long-term public key. We define a

new algorithm `traderAccountScript` which will be used to derive the `pkScript` for a given trader:

<pre> traderAcctScript(P_t, S_t, P_A, B_{k_i}) 1 : $t \leftarrow H(B_{k_i} S_t P_t)$ 2 : $P'_t \leftarrow P_t + (t * G)$ 3 : $P'_A \leftarrow P_A + H(P'_t P_A) * G$ 4 : witnessScript $\leftarrow \{$ 5 : P'_t OP_CHECKSIGVERIFY 6 : P'_A OP_CHECKSIG 7 : OP_IFDUP 8 : OP_NOTIF 9 : T_{blocks} OP_CHECKLOCKTIMEVERIFY 10 : OP_ENDIF 11 : $\}$ 12 : return witnessScript </pre>

The above script can either be spent via the time out clause using the following witness stack, using `nil` value passes an empty signature to force execution of the timeout clause.

`nil traderSig witnessScript`

Or via the normal spend path way which will be used to authorized batches, account closing, and any other account modifications:

`auctioneerSig traderSig witnessScript`

Note that each trader starts using the current batch key at the time they joined the marketplace. However, as we'll see below in the execution section, a trader's key gets rotated with each batch they participate in, meaning that the set of batch keys used within a trader's account output script will eventually become de-synchronized across the market unless all trader's participate in all batches, which is unlikely.

Given the `traderAccountScript` algorithm, we'll now define our implementation of the set of `Account Operations` methods:

NewAccount ($1^\lambda, P_{\text{auction}_p}$)
<pre> 1: $k_t \leftarrow \mathbb{Z}_p$ 2: $P_t \leftarrow k_t * G$ 3: acctScript \leftarrow traderAcctScript(P_t, S_t, P_A, B_{k_i}) 4: accountTxn \leftarrow makeTxn(acctScript) 5: traderSignTxn(accountTxn) 6: broadcastTxn(accountTxn) 7: $\Psi \leftarrow$ await confirmation(accountTxn) 8: return Ψ </pre>
ModifyAccount ($\Psi, P_{\text{auction}_p}$)
<pre> 1: if inPendingBatch(Ψ) 2: return (0, Ψ) 3: 4: $B_{k_i+1} \leftarrow$ IncrementBatchKey(Ψ, B_{k_i}) 5: acctScript \leftarrow traderAcctScript(P_t, S_t, P_A, B_{k_i+1}) 6: 7: $txInputs \leftarrow$ readInput clientStream 8: $txOutputs \leftarrow$ readInput clientStream 9: action = \perp 10: match : 11: acctScript not in $txOutputs$: 12: action = CLOSE 13: $txOutputs.acctOutput.Value > \Psi.Value$: 14: action = DEPOSIT 15: $txOutputs.acctOutput.Value < \Psi.Value$: 16: action = WITHDRAW 17: 18: accountTxn \leftarrow makeTxn(acctScript, action, $txInputs$, $txOutputs$) 19: traderSignTxn(accountTxn) 20: broadcastTxn(accountTxn) 21: 22: $\Psi' \leftarrow$ await confirmation(accountTxn) 23: return Ψ' </pre>

Notice how the **batchKey** is incremented for all account modification operations. In this manner, the **batchKey** also serves as a sequence number within the scripts to ensure that no scripts are re-used within the system.

Users are able to recover their accounts themselves if data has been lost by scanning the chain for the above instances of **newAcctScript** and utilizing the timeout clause within the account script of a trader's account.

7.1.3 Channel Leases in the Lightning Network

Now that we have our concrete account structure, we'll move on to presenting a concrete instantiation of a Channel Lease based on today's widely used payment channel implementation within Lightning Network.

Channel Lease Duration Enforcement

The unique component that sets apart a regular channel from one that was created via a channel lease contract is *duration* enforcement. As a channel lease contracts states the capital must be committed to the network for a minimum period of time, in order to implement this in a trust-minimized manner, we must enhance the existing channel format (cite) used in the Lightning Network today. Our tool of choice for creating minimum duration enforced channels is the `OP_CHECKLOCKTIMEVERIFY` (cite) op-code. Minimally, we need to enforce the following conditions:

- The **Liquidity Maker** involved in a channel lease cannot sweep the funds in their settled commitment output as manifested on their commitment or the commitment of the remote party until D_{block} Bitcoin blocks has passed since the creation of the channel lease.
- The **Liquidity Maker** also cannot fully sweep any funds that are suspended within HTLC outputs until D_{block} Bitcoin blocks has passed.

The second item is of great importance to ensure that the seller of a channel lease can't just move all the committed funds into HTLCs, then close the channel and be fully refunded, netting the lease premium in the process. Instead, we need to ensure that the node is able to resolve any HTLCs on-chain if needed, while still being forced to commit the funds in ancestors of the multi-sig funding output, until the lease duration has expired.

We now make a small modification first to the settled local output of the **Liquidity Maker**:

```
OP_IF
  <revoke key>
OP_ELSE
  <lease duration in blocks>
  OP_CHECKLOCKTIMEVERIFY
  OP_DROP

  <delay in blocks>
  OP_CHECKSEQUENCEVERIFY
  OP_DROP

  <delay key>
OP_ENDIF

OP_CHECKSIG
```

Next, we make a similar modification to the settled remote output of the **Liquidity Maker**, assuming anchor output based channels (cite) are used:

```
<localKey> OP_CHECKSIGVERIFY
<lease duration> OP_CHECKLOCKTIMEVERIFY
1 OP_CHECKSEQUENCEVERIFY
```

Finally, we modify the offered HTLC outputs of the **Liquidity Maker** for their local commitment transaction:

```
OP_IF
  <revoke key>
OP_ELSE
  <lease duration in blocks>
  OP_CHECKLOCKTIMEVERIFY
  OP_DROP

  <delay in blocks>
  OP_CHECKSEQUENCEVERIFY
  OP_DROP

  <delay key>
OP_ENDIF

OP_CHECKSIG
```

A less trust-minimized version of channel lease duration enforcement is possible simply by having the **Liquidity Taker** refuse a cooperative channel closure until the lease duration has expired. With this non-script based enforcement, the only direct option the **Liquidity Maker** has is to *force-close* their channel. However, the operator of an auction venue can request additional information along-side orders before batch execution (as detailed below) to identify premature force closes on-chain in order to penalize the renegading party.

Channel Lease Funding Protocol

Given the existence of a new modified channel C_L that supports the channel lease protocol, we require a new way to fund the channels as a channel lease itself will *partially-bind* the following parameters of a given channel:

- The size of the channel itself.
- The two multi-sig scripts used within the channel.
- The duration of the channel lease itself.

Rather than modifying the base funding protocol of the Lightning Network, we've opted to instead provide a *new* "side-loadable" partial funding binding API we call a channel shim (cite). A channel shim allows a party A to *expect* certain parameters of the channel itself. Given the `pendingChanID` which is used to identify an unfinalized channel within the Lightning Network (cite), an algorithm `RegisterChannelShim(Γ, P_a , outPoint)` instructs a Lightning Node to use a *prearranged* public key as specified within channel lease Γ for its portion of the multi-sig output used within the ultimate channel.

The existence of the channel shim API allows a node to proceed in the set up of a channel for which it may not yet know the *full* funding transaction to. Given the lease, and the expected outpoint, both sides can fully sign the commitment transaction in a manner that doesn't incur any risk as due to the account structure, they'll need to sign off on the batch funding transaction itself before it can be broadcast.

7.1.4 Order Structure

With our concrete account structure and channel lease semantics in place, we'll now outline the precise structure of orders as implemented in **Lightning Pool**. As a CLM is a *sealed-bid* auction, the set of active orders within an auction epoch isn't known to a given participant within the auction. Instead, bids are submitted directly to the auctioneer, and may optionally be cancelled between batch epochs as well. Using the fundamental **unit** notion for expressing the *quantity* of an order, we permit partial matching in addition to specifying a *minimum* matchable amount by adding a new field M_{match} to an order within the Δ_{aux} set of additional order attributes.

Order Tag Generation & Validation

Next, we specifying order tag generation and validation. Before verifying an order, the order Θ itself is serialized in order to generate the message digest of the order. This is done by concatenating each item of the order into a single byte stream, with the set of Δ_{aux} attributes specifying a custom serialization mechanism.

For our tag generation, we opt to utilize an SUF-CMA signature scheme. As we target the base Bitcoin blockchain, **Lightning Pool** utilizes Schnorr signatures implemented over the **secp256k1** elliptic curve. In addition to having the trader that submits an order sign the order message digest, we also require that all the backing Lightning nodes of the order also include a signature as well. As the number of backing nodes for a given account may be numerous, rather than accepting individual signatures for each node, we instead require the account operator and all backing nodes to present a *single* Schnorr signature. As we also want have our order tagging scheme be secure against key rogue-key attacks (cite), we select the **MuSig** multi-signature scheme.

Given the **MuSig** multi-signature scheme, we define the following algorithms used in our order tagging scheme:

- Let $\text{Sign}(\{P_0, \dots, P_i\}, M) \rightarrow \sigma$, be an algorithm that returns a valid **MuSig** multi-signature signed by the set of public keys on message M .
- Let $\text{Verify}(\{P_0, \dots, P_i\}, \sigma, M) \rightarrow b$ be an algorithm that returns $b = 1$ if the passed signature is a valid **MuSig** multi-signature signed by the set of public keys for the message M .

Given these algorithms, we now define our order tag generation and validation implementations:

$\text{GenOrderTag}(P_{acct}, \Theta)$ <hr/> $m \leftarrow K_{nonce} \ V_{ver} \ P_{acct} \ \Delta_{base} \ \Delta_{aux}$ $tag \leftarrow \text{Sign}(\{\Theta.L_{pub}, \Theta.M_{pub} \dots\}, m)$ return tag
$\text{VerifyOrderTag}(P_{acct}, \Theta)$ <hr/> $m \leftarrow K_{nonce} \ V_{ver} \ P_{acct} \ \Delta_{base} \ \Delta_{aux}$ $b \leftarrow \text{Verify}(\{P_{acct}, P_{M_{pub_0}}, \dots, P_{M_{pub_i}}\}, m)$ return b

We omit the implementations of **SubmitOrder** and **CancelOrder** as they depend on the specific environment in which the auctioneer is implemented in. We only add that the pre-image to an order nonce K'_{nonce} is known only to the agent that places the order. As we touch on within the future direction section, this commitment structure also has a number of independent uses outside of order cancellation.

7.1.5 Node Rating Agencies

Recalling the set of initial requirements that were set out in section 4, it's critical that Lightning Pool reduces information symmetry for the buyer by allowing them to gauge the quality of the node selling a channel lease before they enter into an agreement. In order to achieve this, we introduce the concept of a Node Rating Agency. A rating's agency will allow the buyer of a lease to either query the agency in an ad-hoc manner, or specify that they only wish to be matched with nodes that reside on a certain *tier*.

Let $T_{node} = \{\hat{t}_0, \dots, \hat{t}_n\}$ be the set of possible ratings that can be given to a Lightning Node with \hat{t}_0 containing *all* known Lightning Nodes, and $|\hat{t}_i| > |\hat{t}_{i+1}| \forall i < n$, where n is the number of available tiers. In other words, we create a series of node sets, with the higher node tiers having less nodes than lower tiers. This creates a natural system of concentric circles where as the tier index increases, the set of nodes shrinks, and eventually only higher quality nodes remain.

Building off this notion of tiered node sets, we define the following algorithm for our node rating agency:

$\text{NodeTier}(L_{pub})$ <hr/> $\text{nodeMap} \leftarrow \text{map}(T_{node})$ $t_{node} = \text{nodeMap}[L_{pub}]$ return t_{node}

Given this algorithm, we now specify our of the additional auxiliary order attribute Δ_{aux} as:

$$(N_{tier}, \dots) = \Delta_{aux}$$

During match making, we then require the constraint that a given bid order Θ_{bid} will only be matched with an ask order Θ_{ask} if the following constraint is

met:

$$\Theta_{bid}.N_{tier} \geq \text{NodeTier}(\Theta_{ask}.L_{pub})$$

7.1.6 Uniform Price Market Clearing & Matching

With our order structure, and the notion of the Node Rating Agency outlined, we now move on to the concrete market clearing and match making within Lightning Pool.

Order Matching

As mentioned, due to the set of additional constraints outside of simply the posted price and available supply, we employ a multi-attribute match making algorithm. Specifically we opt to utilize a greedy algorithm for the purpose of match making, rather than attempt to find an optimal solution using techniques such as mixed integer linear programming. In this section, we focus primarily on the set of base attributes, leaving the consideration of a trader's auxiliary attribute preferences to later work.

First, we define an abstract algorithm which will be used to determine if a given ask order Θ_{ask} is compatible preference-wise to a given bid order Θ_{bid} :

- **MatchPossible**($\Theta_{bid}, \Theta_{ask}$) $\rightarrow (b, n_{unit})$. This algorithm returns $b = 1$ if the given ask and bid are compatible from a match making perspective. If the orders are compatible, then n_{unit} represents the number of units that can be matched across the two orders.

Given this function we define our implementation of the **MatchMake** algorithm:

MatchMake ($O: \{\Theta_0, \dots, \Theta_n\}$)
1: matchSet $\leftarrow \{\}$ 2: asks $\leftarrow \text{sort}(\text{filter}(O, \Theta_i.O_{type} == \text{Ask}))$ 3: bids $\leftarrow \text{sort}(\text{filter}(O, \Theta_i.O_{type} == \text{Bid}))$ 4: for <i>bid</i> in <i>bids</i> : 5: for <i>ask</i> in <i>asks</i> : 6: if MatchPossible (<i>bid</i> , <i>ask</i>) : 7: matchSet = matchSet \cup (<i>ask</i> , <i>bid</i>) 8: return matchSet

We note that several optimizations here are possible to reduce the worst-case running time of the algorithm which we leave open for future work. We also assume that the set of valid orders has been filtered out before being passed into this algorithm based on the current target batch fee rate and the posted max batch fee rate of each order.

Uniform Price Clearing

Once we've had our set of candidate matches, we'll now move on to the market clearing phase. During the market clearing phase, two distinct operations are carried out:

- A trader's account state is updated to reflect any lease premiums earned due to matches, chain fees paid in the batch execution transaction, fees paid to the auctioneer, and finally the debit for any sold channels from their account.
- We determine the uniform clearing price for a given potential batch.

These two actions comprise the **ClearMarket** and **MarketClearingPrice** algorithms. For our market clearing price, we select the Last Accepted Bid market clearing rule (cite), choosing to go with a buery's mid marking price. Given this price clearing algorithm, we now define our implementaion of the market clearing algorithms:

MarketClearingPrice (Φ_b) <hr/> 1: lastPair $\leftarrow \Phi_b[\text{len}(\Phi_b) - 1]$ 2: return lastPair .bid. α_{rate}
ClearMarket ($\Psi_A, \Phi_b, \{\Psi_0, \dots, \Psi_n\}, c_{price}$) <hr/> 1: leases $\leftarrow \{\}$ 2: accts $\leftarrow \{\}$ 3: for orderPair in Φ_b : 4: lease $\leftarrow \text{newLease}(\text{orderPair})$ 5: orderPair .taker.balance $\mathrel{-=}$ lease .premium 6: orderPair .maker.balance $\mathrel{+=}$ lease .premium 7: 8: orderPair .taker.balance $\mathrel{-=}$ exeFee (lease .amt) 9: orderPair .maker.balance $\mathrel{-=}$ exeFee (lease .amt) 10: Ψ_A .balance $\mathrel{+=}$ exeFee (lease .amt) 11: 12: leases $\leftarrow \text{leases} \cup \text{lease}$ 13: accts [orderPair .taker] \leftarrow orderPair .taker 14: accts [orderPair .maker] \leftarrow orderPair .maker 15: 16: return (Ψ_A , leases , accts)

Notice that we omit the observance of chain fees, as that will be applied to each input/output during the later batch construction phase.

7.1.7 The Batch Execution Transaction

Once we've cleared the market, we'll now move onto the batch execution phase. During this phase, we'll construct the batch transaction which executes a given batch, and also gather all the necessary witnesses for each participant of the batch so we can properly spend their on-chain account outputs. Remember that due to the non-custodial structure, a trader will fully validate a given batch before they sign off on it.

Batch Transaction Construction

A given batch transaction contains the following inputs: the set of trader account inputs involved in the batch, and the input of the master auctioneer itself. In addition to these inputs which can only be spent each each user authorized the proposed batch, we also add the following outputs: a trader's new incremented account output which reflects the market clearing, the set of channels created as part of the channel lease, and the incremented auctioneer account. Note that the format of the batch transaction itself may change multiple times during this phase if participants reject the batch, or if fee changes causing the auctioneer to consider a subset of the prior set of orders.

Taking this into consideration, we define our implementation of **ConstructBatch** as follows:

ConstructBatch(Δ_i)
1 : $tx \leftarrow \text{newTx}()$
2 : $tx.\text{addIn}(\Delta_i.\Psi_A)$
3 :
4 : for $acct$ in $\Delta_i.\Psi$:
5 : $tx.\text{addIn}(acct.\text{prevOut})$
6 : for $acct$ in $\Delta_i.\Psi'$:
7 : $acct.\text{Value} -= \text{feeShare}(\Delta_i, acct)$
8 : $tx.\text{addOut}(acct.\text{txOut})$
9 :
10 : for $lease$ in $\Delta_i.\Gamma$: $tx.\text{addOut}(lease.\text{txOut})$
11 :
12 : $\Psi_A.\text{value} -= \text{feeShare}(\Delta_i, \Psi_A)$
13 : $tx.\text{addOut}(\Delta_i.\Psi'_A)$
14 : return tx

Batch Transaction Execution

Once the batch has been constructed, the auctioneer then needs to propose the batch to each trader, and collect the necessary set of signatures required to spend each trader's account input. During this execution phase, we assume that this is the final set of traders that wish to be a part of this batch. Once the auctioneer has all the necessary signatures to broadcast a batch, the batch execution transaction can be broadcast, ending this auction epoch.

We define batch execution in the context of Lightning Pool as follows:

ExecuteBatch (B_{t_i})
1 : $b \leftarrow 0$ 2 : for $i, input$ in $B_{t_i}.inputs$: 3 : $witness \leftarrow \text{awaitreqSig}(input)$ 4 : $B_{t_i}.witness \leftarrow witness$ 5 : 6 : if $\text{scriptVerify}(B_{t_i})$: 7 : return 1 8 : else 9 : return 0

7.2 The Lightning Pool Shadowchain

In this section, we complete the Lightning Pool system by demonstrating out its implemetnation of a Channel LEase Marketplace can be impeneted using our shadowchain application overlay framework.

7.2.1 Lightning Pool Accounts as Lifted UTXOs

First, we link the concept of our non-custodial accounts in the CLM realm to a lifted UTXO. The process of lifting and unlifting a UTXO is simply a series of operations required to create, modify or close an account:

LiftUTXO ($T_{expiry}, \{U_{N_0}, \dots, U_{N_n}\}, P_0$)
1 : $inputs \leftarrow \{U_{N_0}, \dots, U_{N_n}\}$ 2 : return $\text{NewAccount}(T_{expiry}, P_0, inputs)$
UnliftUTXO (Ψ_U)
1 : $(b,_) \leftarrow \text{ModifyAccount}(\Psi_U, P_{auction_p})$ 2 : return b
ExitChain (ϕ_U, B_{height})
1 : if $\phi_U.T_{blocks} < B_{height}$: 2 : return 0 3 : $(b,_) \leftarrow \text{ModifyAccount}(\Psi_U, P_{auction_p})$ 4 : return b

7.2.2 Auction Batch Proposal

Next, we move unto patch proposal and acceptance. An auction batch within a CLM maps 1:1 to the concept of blocks in the shadowchain domain. Given this insight, we now dfine the **ConstructBlock** and **ProposeBlock** algorithms:

ConstructBlock ($\phi_{live}, T_{xn}, E_{exe}, \Delta_F$) \rightarrow	ProposeBlock (B_S, ϕ_{live})
$\Phi_b \leftarrow \text{MatchMake}(T_{xn})$	$b \leftarrow \text{ValidateBatch}(B_S, \phi_{live})$
$c_{price} \leftarrow \text{MarketClearingPrice}(\Phi_b)$	return b
ClearMarket $\leftarrow \Delta_F$	
$\Delta_i \leftarrow \text{ClearMarket}(\Psi_A, \Phi_b, \Psi, c_{price})$	
return Δ_i	

7.2.3 Shadowchain Batch Execution

Now that we're able to lift/unlift UTXOs and propse blocks within our shadowchain, we now define the series of methods that will be utilized to allow clients to execute the their local version of the state transition function to accept new proposed batches:

CommitBlock (B_S)
1 : $(b, TX_{id}) \leftarrow \text{ExecuteBatch}(B_S)$
2 : return b

7.2.4 Unconfirmed Batch Cut-Through

Recall that a shadowchain block can also be optimistically aggregated into a single block. In the domain of the Lightning Pool shadowchain, combining blocks requires ensuring that all produced channel leases will still exist in the final combined blocks, and the end state of each accountn refelects any consequeutive market clearing opportunities:

CoalesceBlocks ($F: \{B_{S_0}, \dots, B_{S_N}\}$)
1 : inputs $\leftarrow \text{set}(F)$
2 : leases $\leftarrow \text{extractLeases}(F)$
3 : accts $\leftarrow \text{endAcctState}(F)$
4 :
5 : tx $\leftarrow \text{ConstructBatch}(\Delta_i(\text{inputs}, \text{leases}, \text{accts}))$
6 : $b \leftarrow \text{ExecuteBatch}(t)$
7 : return b

7.2.5 Auction Upgrades

Finally, we define the process by which we upgrade the CLM shadowchain itself. As this is an off-chain process, each participant of the shadowchain is able to call the followin algorithm to update in a de-synchronized manner:

UpgradeChain ($\Delta_{new}, E'_{exe}, T'_{xn}$)
1 : return 0

8 Security Analysis

Similar to Loop, the Pool backend server is closed source, but clients are able to fully verify and audit each phase of the auction. At a high level, Pool can be seen as a “shadow chain” anchored in the base Bitcoin blockchain. The shadow chain validates modifications to a subset of the UTXO set (the Pool accounts) with the auctioneer acting as a block proposer to extend the chain. State transitions are validated and accepted by those that are involved in a new block (the auction batch). Newer clients are even able to audit the prior history of the system in order to ensure proper operation. Pool uses the Bitcoin blockchain for what it’s best for: global censorship resistant batch execution.

Leveraging this shadow chain structure, users remain in control of their funds at all times, and will only enter into agreements that they’re able to fully verify, ensuring that channels are properly constructed and that the market is operating as expected. Compared to existing centralized exchanges with off-chain order execution, Pool has a number of attractive security properties:

- As a non-custodial system, users are in control of their funds at all times.
- A purchased LCL will result in the creation of a channel with parameters that capture the preferences expressed in the initial order.
- If the auctioneer server is hacked, the breach doesn’t unilaterally compromise user funds.
- Orders by one trader cannot be used to spoof orders by another trader.
- Clients are able to verify and audit all operations carried out by the server during batch construction including proper order matching.

9 Future Directions

As is clients verify each Shadowchain blocks within the CLM system, but they have no assurance that their order was actually included in the match making function. In this manner, the auctioneer can silently ignore a set of orders. To remedy this, it is possible for the auctioneer to publish an order transparency authenticated data structure to give users a merkle leaf receipt of proper order inclusion.

Rather than sending over a full block, the auctioneer can instead send over a ZKP (cite) of proper block validity. This would improve the privacy of the system.

As is, the maker receives their premium all at once. However it may be possible to set up another uni-directional channel in order to stream the interest in real-time. It may further be possible to allow others to buy/sell the future cash flows of the “coupon channels”.

10 Related Work

11 Conclusion

In this work, we’ve put forth a new abstraction over capital obligations in the Lightning Network, which we call channel leases. Channel leases can be bought and sold on a Channel Liquidity Marketplace directly solving a series of bootstrapping challenges the network faces, commonly referred to as: the inbound liquidity problem. To implement a Channel Lease Marketplace in a secure manner, we put forth the concept of a shadowchain application framework which is of independent use. We then concretely construct Lightning Pool, the first CLM implemented on top of the base Bitcoin blockchain. Lightning Pool allows those with idle capital to earn yield on their Bitcoin, and also allows those that need inbound to receive over the network to obtain a reliable source of incoming payment bandwidth.

12 Acknowledgments

ryan, justin, etc